


Machine Learning Week 8

Unsupervised Learning

1. Clustering




group / clustering:
find some structures

Training set $\{x^{(1)}, x^{(2)}, \dots, x^{(n)}\}$

Applications

- market segmentation
- social network analysis
- organize computer clusters
- astronomical data analysis

2. K-means Algorithm



cluster centroids. (randomly initiated)

step (1) cluster assignment.

- binary assignment of datasets. depending on proximity.

(2) Move centroid.

- move to "mean" of location in all labelled

re-colour.

[Input] .

① k (number of clusters)

② Training set $\{x^{(1)}, x^{(2)}, \dots, x^{(n)}\}$

$x^{(i)} \in \mathbb{R}^n$ [drop $x_0 = 1$ convention]

[Algorithm]

Randomly initialize K cluster centroids $\mu_k, k=1 \dots K$

do $\{$
 for $i=1 \dots n$.

cluster assignment (

$$c_i := \min_k \|x^{(i)} - \mu_k\|^2$$


for $k=1 \dots K$

move centroid (

$$\mu_k := \text{mean}(\text{pts assigned to cluster } k) \\ \in \mathbb{R}^n$$

3. K-means for non-separated clusters.

T-shirt song



Optimization objective

K-means optimization objective.

- c^i : index.
- μ_k : cluster centroid $k \in \mathbb{R}^n$
- μ_{c^i} = cluster centroid of cluster to which $x^{(i)}$ has been assigned.

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$



[Algorithm]

Randomly initialize K cluster centroids $\mu_k, k=1 \dots K$

do \downarrow
for $i=1 \dots m$.

cluster assignment $\left(\begin{array}{l} c^i := \min_k \|x^{(i)} - \mu_k\|^2 \\ \quad \quad \quad \uparrow \\ \text{for } k=1 \dots K \end{array} \right.$

move centroid $\left(\begin{array}{l} \mu_k := \text{mean}(\text{pts assigned to cluster } k) \\ \in \mathbb{R}^n \end{array} \right.$

Minimize $J(\dots)$
w.r.t. $c^{(1)} \dots c^{(m)}$.
[holding $\mu_1 \dots \mu_K$]

Min. w.r.t. $\mu_1 \dots \mu_K$

Random Initialization

Rules.

- (1) $K < m$
- (2) Randomly pick K training examples
- (3) set $\mu_k = \text{examples}$

• Might have different clustering (local optimum) \rightarrow try different random initialization.

- Implementation.

for $i = 1 \dots 1000$ {

 randomly initialize K -means

 run K -means. get $c^{(1)} \dots c^{(K)}, \mu_1 \dots \mu_K$

 complete cost function (distortion)



$\hookrightarrow J(c^{(1)}, \dots, c^{(K)}, \mu_1 \dots \mu_K)$

}

Pick one with lowest cost J .

Choose K

1. Elbow method:




2. Pre-defined (f-chart size: 2...5).

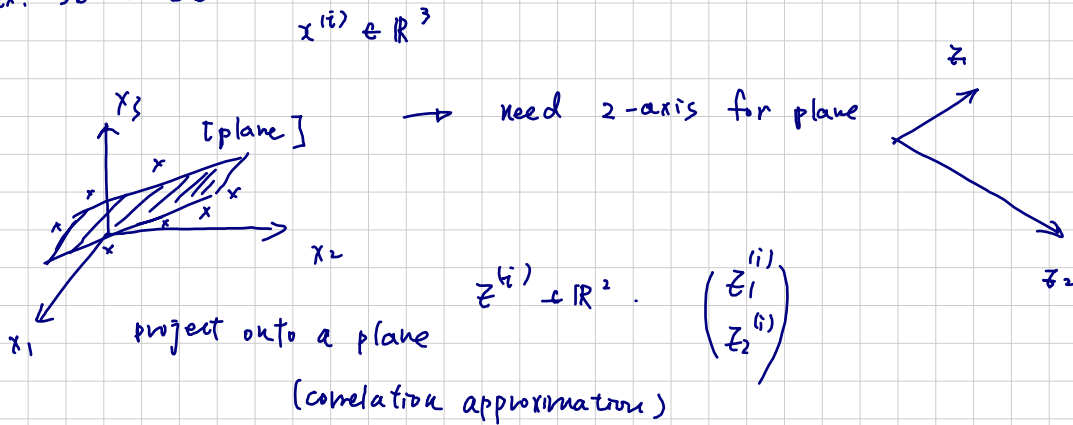
Dimensionality Reduction

1. Motivation I: Data compression

- Redundant data dimension (cm, inch, ...)



ex. $3D \rightarrow 2D$.




procedure. project data to one less dimension. \rightarrow re-coordinate.

2 Motivation II: Data visualization.

- Combine certain features (correlated).

$$x \rightarrow z$$
$$(500) \quad (2)$$

Principal Component Analysis



Try to find a lower dimension surface to
min. projection error.

→ perform mean normali. and feature scaling.

Principal Component Analysis [PCA] problem formulation



2 → 1
- find dir. vector $u^{(1)} \in \mathbb{R}^k$ onto which data projects

$k \rightarrow k$

- find k vectors $u^{(1)}, u^{(2)}, \dots, u^{(k)}$ onto which data projects.

[project data to the linear subset span by $\{u^{(i)}\}_1^k$]

Linear regression



Applies to higher dimension.

Principal component analysis component.

• Data preprocessing

- mean normalisation
- feature scaling

• Reduction $X^{(i)} \in \mathbb{R}^2 \rightarrow Z^{(i)} \in \mathbb{R} \quad [2D \rightarrow 1D]$

• Mathematical derivation (complicated)

Algo.

① Reduce data from \mathbb{R}^n to \mathbb{R}^k

Compute "covariance matrix"

Compute "eigenvectors" of matrix Σ

$$\text{Sigma} \in \mathbb{R}^{n \times n} \quad (n \times 1) \quad (1 \times n)$$

$$\Sigma = \frac{1}{m} \sum_{i=1}^n [X^{(i)}][X^{(i)}]^T = \left(\frac{1}{m}\right) \cdot X^T \cdot X.$$

$$X = \begin{pmatrix} \vdots \\ \vdots \\ \vdots \end{pmatrix}^n$$

→ singular value decomposition

$$[U, S, V] = \text{svd}(\text{Sigma})$$

$$U = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \dots & u^{(m)} \\ | & | & & | \end{bmatrix} \in \mathbb{R}^{n \times n}.$$

first k vectors: $u^{(1)} \dots u^{(k)}$

② Projection.

$$U_{\text{reduce}} = \begin{bmatrix} | & & | \\ u^{(1)} & \dots & u^{(k)} \\ | & & | \end{bmatrix}$$

$$\in \mathbb{R}^{n \times k}$$

$$Z^{(i)} = U_{\text{red}}^T \cdot X^{(i)} = \begin{pmatrix} -u^{(1)}- \\ \vdots \\ -u^{(k)}- \end{pmatrix}^{(k \times n)} X^{(i)}^{(n \times 1)} = (k \times 1)$$

Reconstruction from Compressed Representation

$$z = U_{red}^T x.$$

$$z \in \mathbb{R} \rightarrow x \in \mathbb{R}^2$$

$$x \approx x_{approx} = \underbrace{U_{reduce}}_{(n \times k)} \cdot z. \quad (k \times 1)$$

Choosing k . (# of principal component)

- Average square projection error $\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{approx}^{(i)}\|^2$

- Total variation in data $\frac{1}{m} \sum_{i=1}^m \|x^{(i)}\|^2$

Typically, choose k to be min s.t.

$$\frac{\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{approx}^{(i)}\|^2}{\frac{1}{m} \sum_{i=1}^m \|x^{(i)}\|^2} \leq 0.01 \quad (1\%) \quad [x]$$

"99% variance retained"

Algorithm.

1) try PCA with $k=1$

2) Compute $U_{reduce}, z^{(1)}, z^{(2)}, \dots, z^{(m)}$
 $x_{approx}^{(1)} \dots x_{approx}^{(m)}$

3) check $[x]$

4) $k=k+1 \rightarrow 2$

- $[U, S, V] = \text{svd}(\text{sigma})$

$\left\{ \begin{array}{l} n \times n \text{ diagonal} \end{array} \right.$, $S = \begin{pmatrix} s_{11} & & 0 \\ & \ddots & \\ 0 & & s_{nn} \end{pmatrix}$

For given k , $[x]$ can be computed as $\left(1 - \frac{\sum_{i=1}^k s_{ii}}{\sum_{i=1}^n s_{ii}} \right)$
 > 0.99

Advice for applying PCA

Supervised Learning Speed Up

$$x^{(i)} \in \mathbb{R}^{10000}$$

$$(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$$

Extract inputs:

$$\text{Unlabelled dataset } x^{(1)}, x^{(2)} \dots x^{(m)} \in \mathbb{R}^{10000}$$

$$\downarrow \text{PCA (defined using only the training set)} \\ z^{(1)} \dots z^{(m)} \in \mathbb{R}^{1000}$$

New training set

$$(z^{(1)}, y^{(1)}) \dots (z^{(m)}, y^{(m)}) \xrightarrow{\text{train.}} h_{\theta}(z) = \frac{1}{1 + \exp(-\theta^T z)}$$

Example

$$x \rightarrow z \rightarrow h(z) \quad (\text{prediction})$$

Applications

- Compression \rightarrow choose k by % variance retained.
- Visualization $\rightarrow k=2,3$

What not to do \rightarrow prevent overfitting

- Use $z^{(i)}$ instead of $x^{(i)}$. reduce # of features from n to k .
 \rightarrow fewer features, less likely to overfit. (Why \rightarrow PCA throws away information w/o knowing y ?)
- should use 1.

PCA usage.

Design of ML sys.?

- get training set

- [- run PCA reduce $x^{(i)}$ \rightarrow $z^{(i)}$] do w/o this step first.

- train logistic.

- test set